

# Neural Networks

Slide credits :Vineeth N Balasubramanian




आई आई टी हैदराबाद  
IIT Hyderabad

# Classification Methods

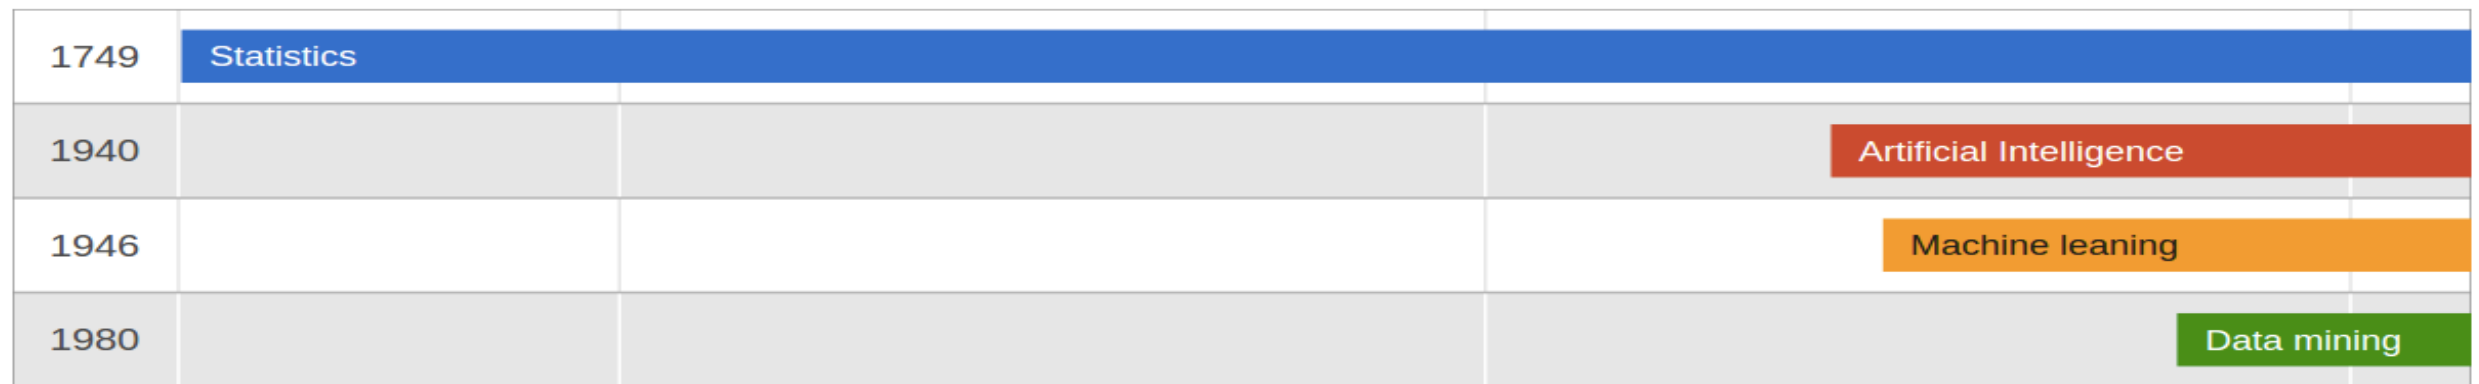
- k-Nearest Neighbors
- Decision Trees
- Naïve Bayes
- Support Vector Machines
- Logistic Regression
- Neural Networks (Deep Learning)
- Ensemble Methods (Boosting, Random Forests)

# Neural Networks (aka) Deep Learning

## Introduction




Deep learning: A sub-area of machine learning, that is today understood as representation learning




# Deep Learning

## Introduction

- Rebirth of neural networks
- Inspired by the human brain (networks of neurons)



# History of Deep Learning



McCulloch-Pitts



Rosenblatt



Widrow-Hoff



Minsky-Papert



Rumelhart-Hinton-Williams






LeCun



Hinton-Ruslan




Krizhevsky-Sutskever-Hinton



# Deep Learning


How do they learn?



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?




Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*A dataset*

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	



Source: Prof Corne. Heriot-Watt University, UK



# Deep Learning

How do they learn?

*Training data*

<i>Fields</i>	<i>class</i>
---------------	--------------

1.4 2.7 1.9	0
-------------	---


3.8 3.4 3.2	0
-------------	---

6.4 2.8 1.7	1
-------------	---

4.1 0.1 0.2	0
-------------	---

etc ...

Initialise with random weights



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Present a training pattern**



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Present a training pattern**



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK



# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9              0


3.8 3.4 3.2              0

6.4 2.8 1.7              1

4.1 0.1 0.2              0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

Source: Prof Corne, Heriot-Watt University, UK

*Training data*

*Fields*                      *class*

1.4 2.7 1.9              0

3.8 3.4 3.2              0


6.4 2.8 1.7              1

4.1 0.1 0.2              0

etc ...

Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments, reduce the error


And so on ....



Called “Gradient Descent”


# The decision boundary perspective...

Initial random weights




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**




# The decision boundary perspective...

**Present a training instance / adjust the weights**



# The decision boundary perspective...


Eventually ....





# Neural Networks Training: Backpropagation

## Perceptrons




$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq \theta \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < \theta \end{cases}$$

- Learn weights such that an objective function is maximized.
- What objective function should we use?
- What learning algorithm should we use?

# Neural Networks Training: Backpropagation

## Perceptrons




$x_1$	$x_2$	$t$
.8	.3	1
.4	.1	0

$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq q \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < q \end{cases}$$

# Neural Networks Training: Backpropagation

## First Training Instance




$x_1$	$x_2$	$t$
.8	.3	1
.4	.1	0

$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq q \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < q \end{cases}$$

# Neural Networks Training: Backpropagation

## Second Training Instance



$x_1$	$x_2$	$t$
.8	.3	1
.4	.1	0

$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq q \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < q \end{cases}$$

$$\Delta w_i = (t - z) * c * x_i$$

# Neural Networks Training: Backpropagation

## Perceptron Rule Learning

$$\Delta w_i = c(t - z) x_i$$

- Where  $w_i$  is the weight from input  $i$  to perceptron node,  $c$  is the learning rate,  $t_j$  is the target for the current instance,  $z$  is the current output, and  $x_i$  is  $i^{\text{th}}$  input
- Least perturbation principle
  - Only change weights if there is an error
  - small  $c$  sufficient to make current pattern correct
  - Scale by  $x_i$
- Create a perceptron node with  $n$  inputs
- Iteratively apply a pattern from the training set and apply the perceptron rule
- Each iteration through the training set is an *epoch*
- Continue training until total training set error ceases to improve
- Perceptron Convergence Theorem: Guaranteed to find a solution in finite time if a solution exists

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Extension of perceptrons to multiple layers
- 1. **Initialize** network with **random** weights
- 2. **For all** training cases (**called examples**):
  - **a.** Present training inputs to network and calculate output
  - **b.** For all layers (starting with output layer, back to input layer):
    - i. Compare **network output** with **correct output** (error function)
    - ii. **Adapt weights** in current layer

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Method for **learning weights** in feed-forward (FF) nets
- Can't use Perceptron Learning Rule
  - no **teacher values** are possible for **hidden units**
- Use **gradient descent** to minimize the error
  - **propagate deltas** to **adjust for errors**
  - **backward from outputs** to hidden layers **to inputs**

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- The idea of the algorithm can be summarized as follows :
  1. Computes the **error term for the output units** using the observed error.
  - 2. From output layer, repeat
    - propagating the error term back to the previous layer and **updating the weights between the two layers** until the earliest hidden layer is reached.



# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Initialize weights (typically random!)
- Keep doing epochs
  - **For each** example **e** in training set do
    - **forward pass** to compute
      - $y = \text{neural-net-output}(\text{network}, e)$
      - $\text{miss} = (T - y)$  at each output unit
    - **backward pass** to calculate deltas to weights
    - update all weights
  - end
- until **tuning set error stops improving**


Forward pass

Backward pass explained in next few slides

# Neural Networks Training: Backpropagation

## Gradient Descent

- Think of the  $N$  weights as a point in an  $N$ -dimensional space
- Add a dimension for the observed error
- Try to minimize your position on the “error surface”



# Neural Networks Training: Backpropagation

Compute  
deltas

Gradient Descent

- Trying to make **error decrease the fastest**
- **Compute:**
  - $\text{Grad}_E = [dE/dw_1, dE/dw_2, \dots, dE/dw_n]$
- **Change i-th weight by**
  - $\text{delta}_{w_i} = -\text{alpha} * dE/dw_i$
- We need a **derivative**!
- Activation function must be continuous, differentiable, non-decreasing, and easy to compute



Derivatives of error for weights

# Neural Networks Training: Backpropagation

Updating Hidden-to-Output

$$\min_{\mathbf{W}, \mathbf{v}} \sum_n \frac{1}{2} \left( y_n - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}_n) \right)^2$$

$$\nabla_v = - \sum_n e_n \mathbf{h}_n$$

# Neural Networks Training: Backpropagation

## Updating Hidden

$$\mathcal{L}(\mathbf{W}) = \frac{1}{2} \left( y - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}) \right)^2$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}_i} = \frac{\partial \mathcal{L}}{\partial f_i} \frac{\partial f_i}{\partial \mathbf{w}_i}$$

$$\frac{\partial \mathcal{L}}{\partial f_i} = - \left( y - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}) \right) v_i = -e v_i$$

$$\frac{\partial f_i}{\partial \mathbf{w}_i} = f'(\mathbf{w}_i \cdot \mathbf{x}) \mathbf{x}$$

Here we have general formula with derivative, next we use for sigmoid

- for sigmoid the derivative is,  $f'(x) = f(x) * (1 - f(x))$

$$\nabla_{\mathbf{w}_i} = -e v_i f'(\mathbf{w}_i \cdot \mathbf{x}) \mathbf{x}$$

Derivative of activation function


# Making Choices

## Backpropagation

- Number of hidden layers – *empirically determined*
  - Too few ==> can't learn
  - Too many ==> poor generalization
- Number of neurons in each hidden layer – *empirically determined*
- Activation functions
- Error/loss functions
- Learning rate
- Gradient descent methods


# Activation Functions

Sigmoid




$$y = \frac{1}{1 + e^{-x}}$$

tanh




$$y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

ReLU



$$y = \max(0, x)$$

Leaky ReLU



$$y = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{if } otherwise \end{cases}$$

# Loss Functions

- Euclidean loss / Squared loss  $L = \frac{1}{2} \|x_i - y_i\|_2^2$ 
  - Derivative w.r.t.  $x_i$   $\frac{\partial L}{\partial x_i} = x_i - y_i$
- Soft-max loss/multinomial logistic regression loss

$$p_i = \frac{e^{x_i}}{\sum_k e^{x_k}} \quad L = - \sum_i y_i \log(p_i)$$

- Derivative w.r.t.  $x_i$   $\frac{\partial L}{\partial x_i} = p_i - y_i$
- Also called: Cross-entropy loss

neural networks loss function is non-convex and optimization is sensitive to their initialization.



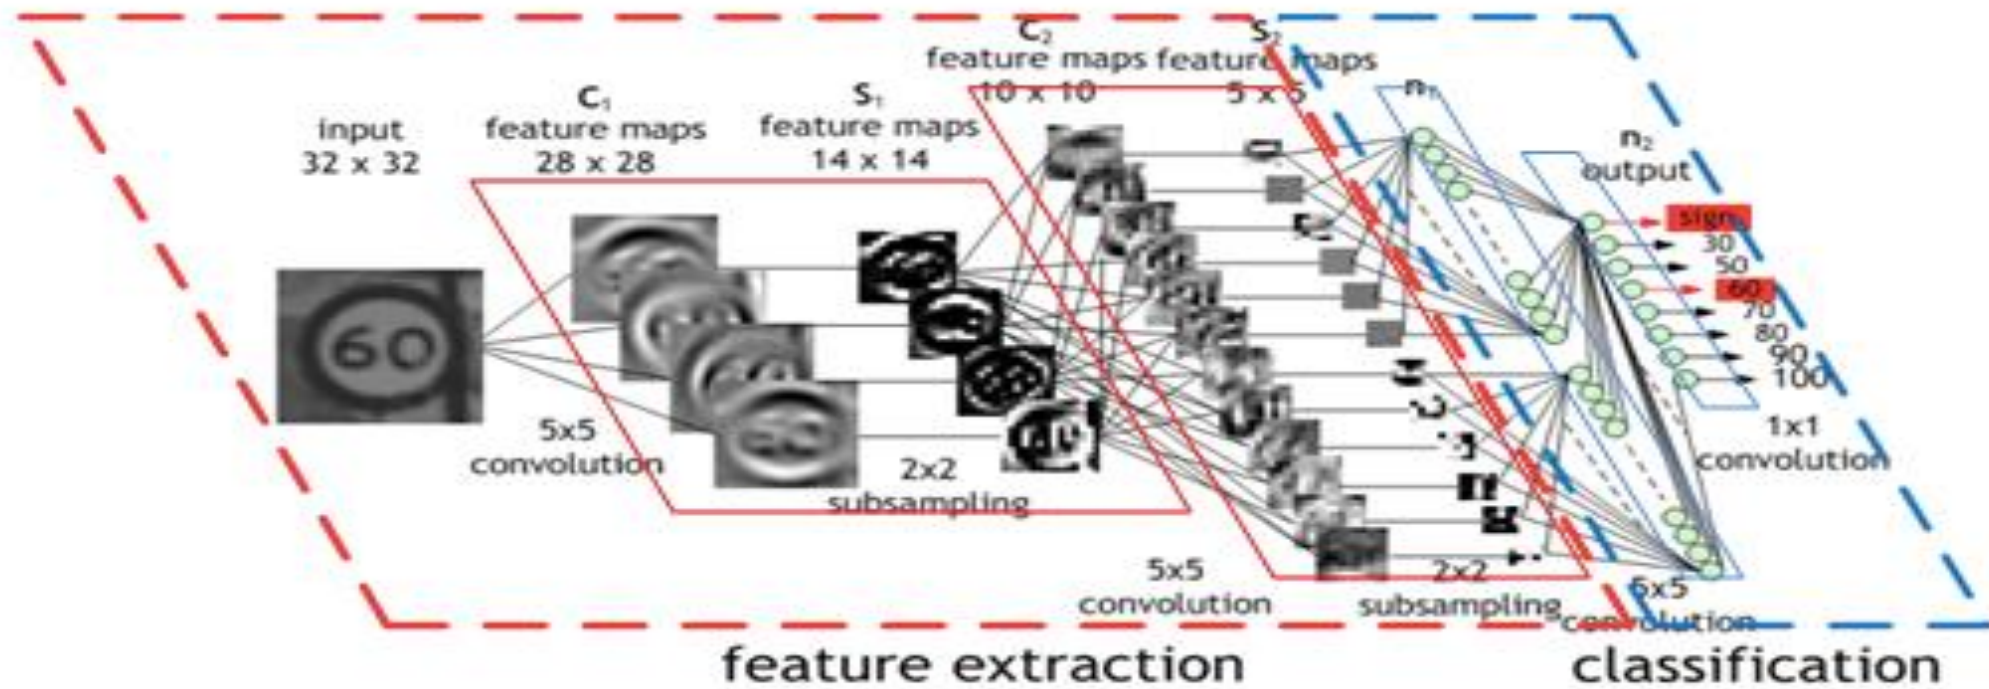
# Gradient Descent Methods

- Batch gradient descent (vs) Stochastic gradient descent (vs) Mini-batch stochastic gradient descent
  - Mini-batch SGD the most popularly used
- Using momentum
- Setting learning rate
  - Fixed learning rate
  - Using learning rate schedules
  - Adaptive learning rate methods: Adam, Adadelata, Adagrad, RMSProp

# Deep Learning Architectures


## Variants

### Convolutional Neural Networks for Image and Video Understanding




# CNNs: Typical Architecture

- A typical deep convolutional network



# Deep Learning in Computer Vision: The Turning Point

## AlexNet in the ImageNet Challenge



## ImageNet Classification with Deep Convolutional Neural Networks

*ImageNet Classification Task:*

*Previous Best: ~25% (CVPR-2011)*

*AlexNet : ~15 % (NIPS-2012)*

Alex Krizhevsky  
University of Toronto  
kriz@cs.utoronto.ca


Ilya Sutskever  
University of Toronto  
ilya@cs.utoronto.ca

Geoffrey E. Hinton  
University of Toronto  
hinton@cs.utoronto.ca

# Deep Learning Architectures

## Variants


### Recurrent Neural Networks for Time Series and Sequence Data Understanding



# Deep Learning Architectures

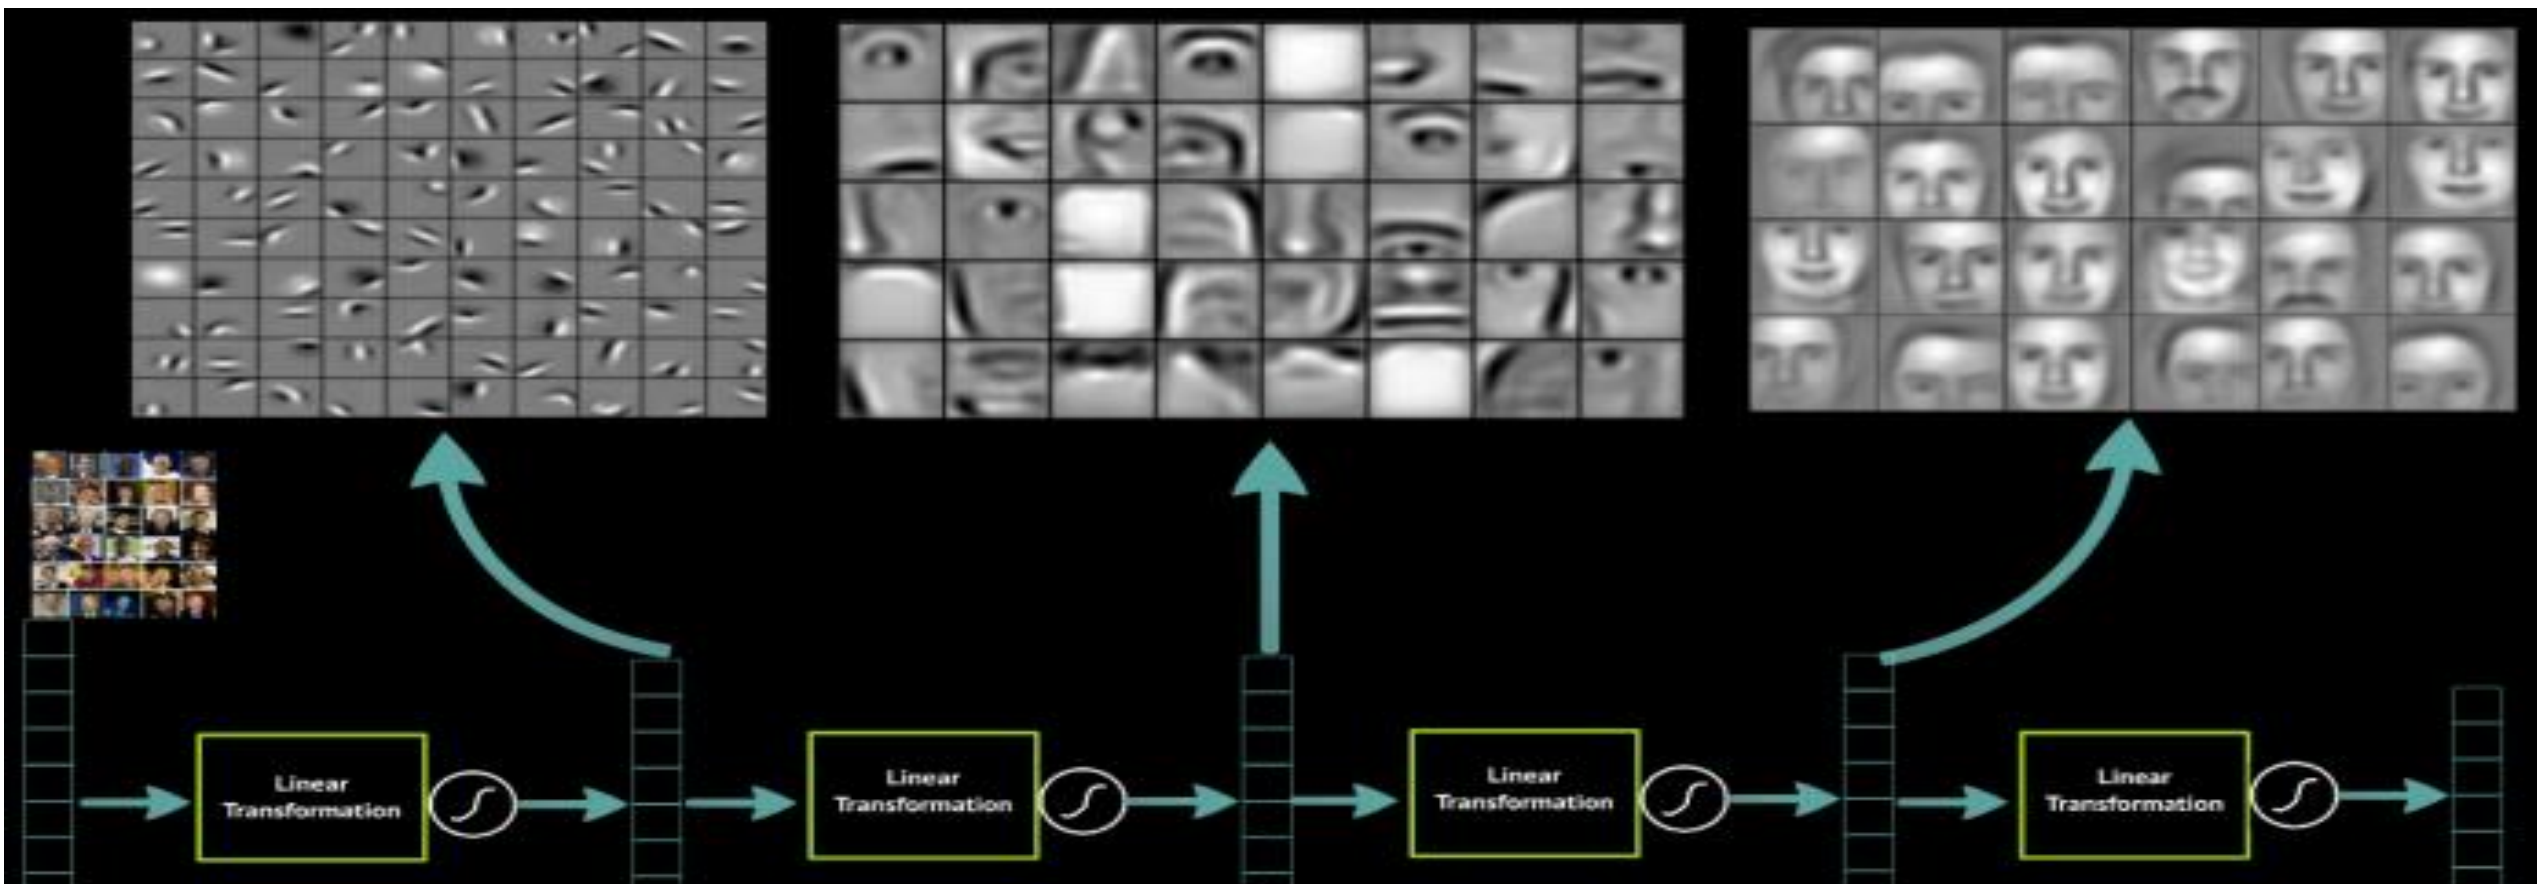
## Variants

### Deep Autoencoders for Dimensionality Reduction



# Deep Learning

Why is it successful?



# Deep Learning

Why has use of multiple layers led to success?


- Captures compositionality: world is compositional!
  - **Image recognition:** Pixel → edge → texture → motif → part → object
  - **Text:** Character → word → word group → clause → sentence → story
  - **Speech:** Sample → spectral band → sound → ... → phone → phoneme → word
- Exploiting compositionality gives an exponential gain in representational power



# Deep Learning

## Why is it successful?

- Learns representations of data that are useful (Other ML algorithms are “shallow”)
- Similar to the human brain
- Then, why was it not successful earlier (in the 90s)?
  - Computational power
  - Data power



# Deep Learning

## Applications and Successes

- AlexNet (Object Recognition): The network that catapulted the success of deep learning in 2012



# Deep Learning


## Applications and Successes

- Speech understanding and natural language processing

Apple Siri



Google Now




Windows Cortana



# Deep Learning

## Applications and Successes

- Game playing: <https://www.youtube.com/watch?v=V1eYniJ0Rnk>



# More?

## Where to look?

- One-stop shop
  - <https://github.com/ChristosChristofidis/awesome-deep-learning>
- Check this out for hours of fun and amazement
  - <http://fastml.com/deep-nets-generating-stuff/>
- Books (on Deep Learning)
  - <http://www.deeplearningbook.org>
  - <http://neuralnetworksanddeeplearning.com/>
- Programming
  - Tensorflow, PyTorch, Theano/Pylearn2, Caffe, Torch, Keras

# Readings

- [“Introduction to Machine Learning” by Ethem Alpaydin](#), Chapters 11.1-11.11
- Bishop, PRML, Sec 5.1-5.3, 5.5