

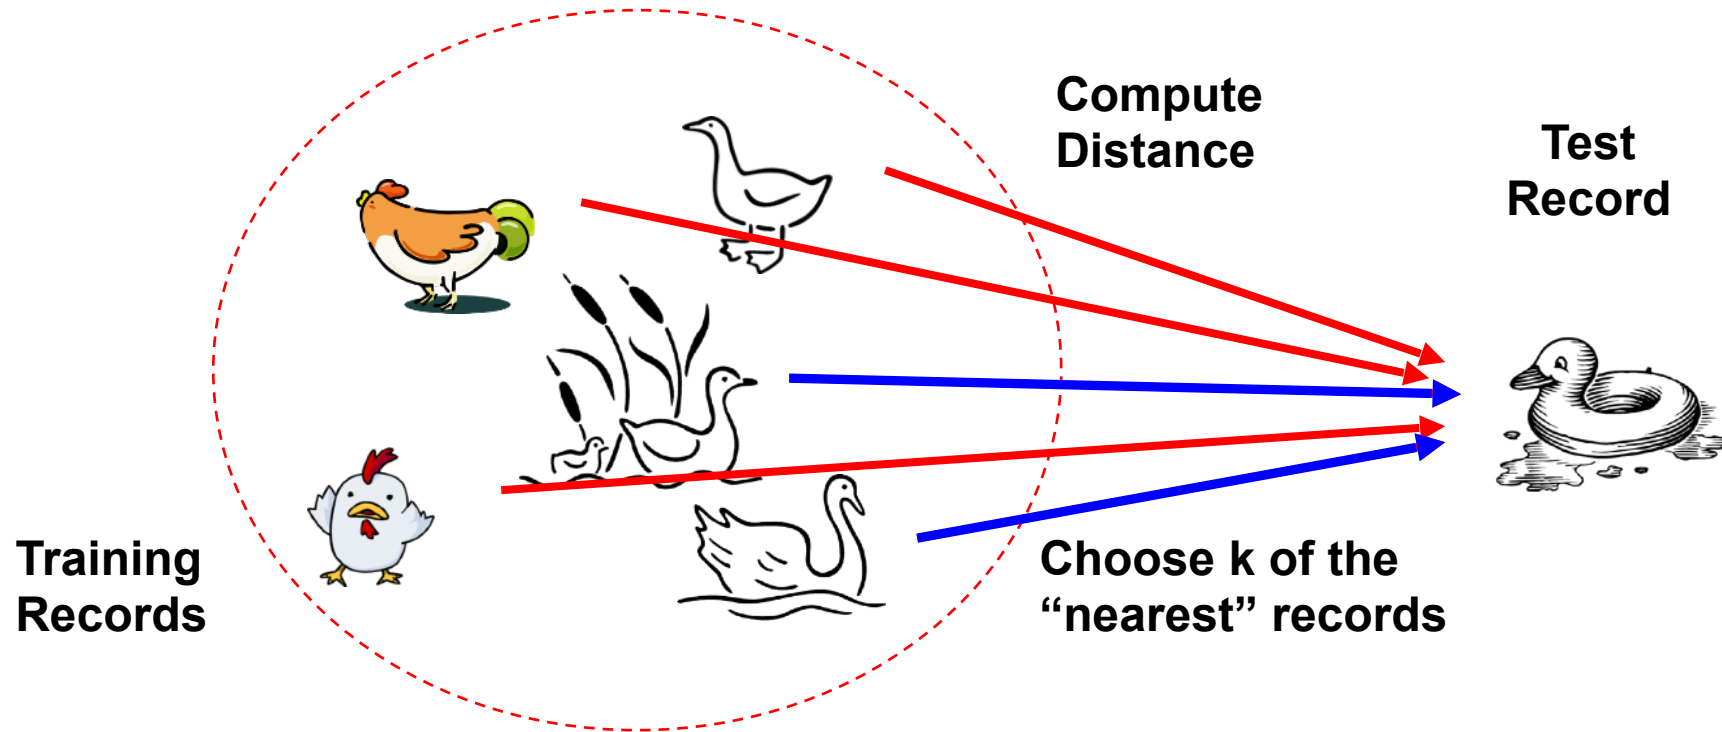
Classifiers: kNN



आई आई टी हैदराबाद
IIT Hyderabad


k-Nearest Neighbors

- Basic idea:
 - If it walks like a duck, quacks like a duck, then it's probably a duck



k-Nearest Neighbors

- Majority vote within the k nearest neighbors



k-Nearest Neighbors

- An arbitrary instance is represented by $(a_1(x), a_2(x), a_3(x), \dots, a_n(x))$
 - $a_i(x)$ denotes features
- Euclidean distance between two instances
 - $d(x_i, x_j) = \sqrt{\sum_{r=1}^n (a_r(x_i) - a_r(x_j))^2}$
- L_p distance
 - $p=2$: Euclidean distance
 - $p=1$: Manhattan distance
 - $p = \infty$: Max distance
 - $p=0$: Count non-zero distance
- In case of continuous-valued target function
 - Mean value of k nearest training examples

Other Distance Metrics


- Cosine Distance Metric $\rho(\vec{x}_1, \vec{x}_2) = \cos(\angle(\vec{x}_1, \vec{x}_2)) = \frac{\vec{x}_1 \cdot \vec{x}_2}{\|\vec{x}_1\|_2 \|\vec{x}_2\|_2}$
- Edit Distance $x_1 = \text{AAATCCCGTAA}$
 $x_2 = \text{AATCGCGTAA}$

Minimum number of insertions, deletions and mutations needed


 $\rho(x_1, x_2) = 2$

k-Nearest Neighbors


- Choosing k is important
 - If k is too small, sensitive to noise points
 - If k is too large, neighborhood may include points from other classes



(a) 1-nearest neighbor



(b) 2-nearest neighbor



(c) 3-nearest neighbor

How to determine k


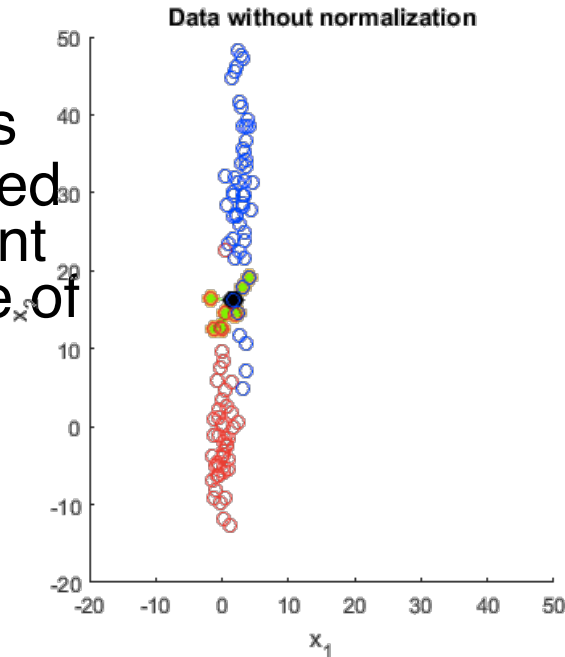
- Determined experimentally (think cross-validation!)
 - Start with $k=1$ and use a test set to validate the error rate of the classifier
 - Repeat with $k=k+2$
 - Choose the value of k for which the error rate is minimum
 - Note: k typically an odd number to avoid ties in binary classification

Pros and Cons

- Pros
 - Highly effective and simple method
 - Trains very fast (“Lazy” learner)
- Cons
 - Curse of dimensionality
 - In higher dimensions, all data points lie on the surface of the unit hypersphere!
 - Closeness in raw measurement space may not be good for the task
 - Storage: all training examples are saved in memory
 - A decision tree or linear classifier is much smaller
 - Slow at query time
 - Can be overcome and presorting and indexing training samples


Improvements

- Distance-Weighted Nearest Neighbors
 - Assign weights to the neighbors based on their 'distance' from the query point (E.g., weight 'may' be inverse square of the distances)
 - Can also learn this -> “**Metric Learning**”
- Scaling (**normalization**) attributes for fair computation of distances




Improvements

- Finding “close” examples in a large training set quickly
 - E.g. Efficient memory indexing using kd-trees
 - In 1-dimension, can reduce complexity from $O(n)$ to $O(\log n)$ – assuming data is sorted
 - Other methods
 - Locality-Sensitive Hashing, Clustering-based methods



Improvements

- Not storing all examples
 - We can label each cell instead and discard the training data



Classification : Evaluation metrics

Accuracy : $\frac{1}{N} \sum_{i=1}^N \mathbb{I}[y_i == \hat{y}_i]$

		Actual Label	
		Positive	Negative
Predicted Label	Positive	True Positive (TP)	False Positive (FP)
	Negative	False Negative (FN)	True Negative (TN)

Accuracy	$(TP + TN) / (TP + TN + FP + FN)$	The percentage of predictions that are correct
Precision	$TP / (TP + FP)$	The percentage of positive predictions that are correct
Sensitivity (Recall)	$TP / (TP + FN)$	The percentage of positive cases that were predicted as positive
Specificity	$TN / (TN + FP)$	The percentage of negative cases that were predicted as negative

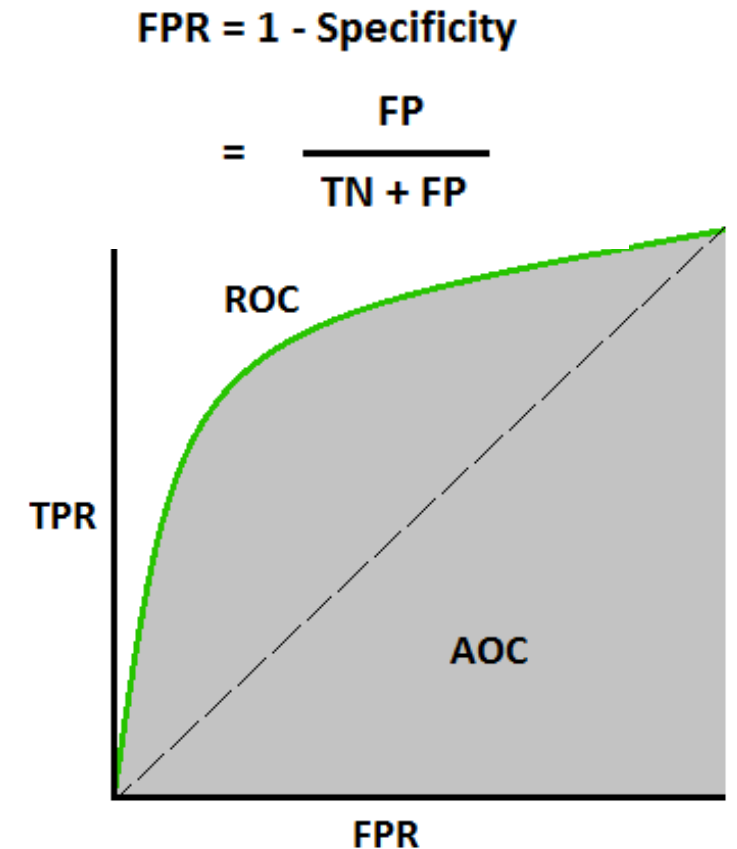
Classification metric : ROC and AUC

- AUC (Area Under The Curve) ROC (Receiver Operating Characteristics) curve.

$$\text{TPR / Recall / Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

- AUC provides an aggregate measure of performance across all possible classification thresholds. One way of interpreting AUC is as the probability that the model ranks a random positive example more highly than a random negative example. Useful for class imbalanced data.



Readings

- Chapters 8, 9, EA Introduction to ML 2nd Edn
- Chapter 14 (Sec 14. 4) + Chapter 2 (Sec 2. 5), Bishop, PRML