


Neural Networks


Deep Learning

Introduction

- Rebirth of neural networks
- Inspired by the human brain (networks of neurons)



History of Deep Learning



McCulloch-Pitts



Rosenblatt



Widrow-Hoff



Minsky-Papert




Rumelhart-Hinton-Williams






LeCun





Hinton-Ruslan



Krizhevsky-Sutskever-Hinton




Neural Networks




Neuron

- Based on how much these incoming neurons are firing, and how “strong” the neural connections are, our main neuron will “decide” how strongly it wants to fire.
- Learning in the brain happens by neurons becoming connected to other neurons, and the strengths of connections adapting over time.
- Receives input from D-many other neurons, one for each input feature. The strength of these inputs are the feature values.




Neural networks

Firing is interpreted as being a positive example and not firing is interpreted as being a negative example



Source: Prof Corne, Heriot-Watt University, UK


Deep Learning



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

Features with zero weight are ignored. Features with positive weights are indicative of positive examples because they cause the activation to increase. Features with negative weights are indicative of negative examples because they cause the activation to decrease.



$$x = -0.06 \times 2.7 + 2.5 \times 8.6 + 1.4 \times 0.002 = 21.34$$

$$a = \left[\sum_{d=1}^D w_d x_d \right] + b$$

Source: Prof Corne, Heriot-Watt University, UK

Perceptron Algorithm

- **Online.** This means that instead of considering the entire data set at the same time, it only ever looks at one example.
- **error driven.** This means that, so long as it is doing well, it doesn't bother updating its parameters

Perceptron Algorithm

Algorithm 5 PERCEPTRONTRAIN(\mathbf{D} , $MaxIter$)

```
1:  $w_d \leftarrow 0$ , for all  $d = 1 \dots D$  // initialize weights
2:  $b \leftarrow 0$  // initialize bias
3: for  $iter = 1 \dots MaxIter$  do
4:   for all  $(x,y) \in \mathbf{D}$  do
5:      $a \leftarrow \sum_{d=1}^D w_d x_d + b$  // compute activation for this example
6:     if  $ya \leq 0$  then
7:        $w_d \leftarrow w_d + yx_d$ , for all  $d = 1 \dots D$  // update weights
8:        $b \leftarrow b + y$  // update bias
9:     end if
10:  end for
11: end for
12: return  $w_0, w_1, \dots, w_D, b$ 
```



Perceptron algorithm

- weight w_d is increased by $y x_d$ and the bias is increased by y .
- goal of the update is to adjust the parameters so that they are “better” for the current example

$$\begin{aligned} a' &= \sum_{d=1}^D w'_d x_d + b' \\ &= \sum_{d=1}^D (w_d + x_d) x_d + (b + 1) \\ &= \sum_{d=1}^D w_d x_d + b + \sum_{d=1}^D x_d x_d + 1 \\ &= a + \sum_{d=1}^D x_d^2 + 1 > a \end{aligned}$$

Perceptron algorithm




- If we make many many passes over the training data, then the algorithm is likely to overfit. On the other hand, going over the data only one time might lead to underfitting
- loop over all the training examples in a constant
- Order is a bad idea. re-permute the examples in each iteration.



Perceptron decision boundary

$$\mathcal{B} = \left\{ x : \sum_d w_d x_d = 0 \right\}$$

normalized weight vectors, $w \cdot x$ is just the distance of x , w is exactly the activation of that example




Deep Learning

How do they learn?

A dataset

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	



Source: Prof Corne. Heriot-Watt University, UK


Deep Learning

How do they learn?

Training data

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	

Initialise with random weights



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Present a training pattern



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Present a training pattern



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Training data

Fields ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

Deep Learning

How do they learn?

Source: Prof Corne, Heriot-Watt University, UK

Training data

Fields *class*

1.4 2.7 1.9 0

3.8 3.4 3.2 0


6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments, reduce the error


And so on



Called “Gradient Descent”


The decision boundary perspective...

Initial random weights




The decision boundary perspective...

Present a training instance / adjust the weights




The decision boundary perspective...

Present a training instance / adjust the weights




The decision boundary perspective...

Present a training instance / adjust the weights



The decision boundary perspective...

Present a training instance / adjust the weights



The decision boundary perspective...

Eventually

